

Ensuring Reproducibility in Machine Learning Experiments: Tools, Best Practices & Real-World Implementation

Timeline: Day 91–105

Focus: Experimentation & Reproducibility

Why Reproducibility Matters in Machine Learning

In machine learning, reproducibility means being able to run the same experiment again and get the same results. It's not just a good habit—it's necessary for others to trust your work, use it, and build on it. If someone can't repeat your results, they might question how reliable your model really is, no matter how impressive it looks.

Over the last two weeks, I worked on making my machine learning experiments reproducible. I used tools like Docker to create consistent environments and DVC (Data Version Control) to track datasets and models. This blog is a reflection of that experience—what tools I used, the problems I ran into, and how I solved them to build a more reliable and sharable workflow.

Topics Covered

Reproducible Research: What It Means

A reproducible ML workflow ensures that:

- Experiments can be re-run with the same results.
- Environments are controlled, so dependency versions don't affect outcomes.
- Data and model versions are tracked, so results can be traced back to exact input settings.

Tools we Used

Docker: Containerization for Environment Management

Docker helped me encapsulate the development environment so that anyone can run the code without worrying about OS or library mismatches.

- Defined all dependencies in a Dockerfile.
- Used docker-compose to manage multiple services (like Jupyter and DVC).

DVC: Data Version Control

DVC made it easier to:

- Track large datasets.
- Version model checkpoints.
- Automate data pipelines with reproducible steps.

Hands-On Tasks

Step 1: Writing Reproducible Experiment Scripts

- Modularized the code using config files (.yaml) for parameters.
- Logged results, metrics, and configurations using MLflow.

Step 2: Integrating with Version Control

- Pushed codebase and DVC config to GitHub.
- Tracked .dvc files for data and model checkpoints.
- Used .gitignore to manage large file tracking separately via DVC.

Step 3: Documenting the Process

I wrote this blog post to:

- Share the overall architecture.
- Detail the setup challenges.
- Provide code snippets for beginners.

Challenges Faced

- DVC Setup with Remote Storage
Took time to configure AWS S3 as remote storage. Had to troubleshoot permission issues.
- Docker Image Size
My first Docker build was over 5GB. I optimized it using slim base images and multi-stage builds.
- Version Drift in Dependencies
Installing dependencies without pinning versions led to inconsistent behavior. Resolved by using requirements.txt with fixed versions and Docker.

Final Deliverables

GitHub Repository

A public repository containing:

- All experiment code

- DVC-tracked data
- Docker setup and usage guide

Research Report & Blog Tutorial

- Best practices guide on reproducibility
- Docker and DVC setup walkthrough
- Sample pipeline with train.py, evaluate.py, and config.yaml

Best Practices for Reproducibility

- Use Docker for all experiments
- Track all data and models using DVC
- Pin all dependency versions
- Log everything: parameters, metrics, and outcomes
- Version control all artifacts — not just code

Conclusion

This project was a key milestone in my machine learning journey, helping me understand the real value of reproducibility in research. By using tools like Docker and DVC, I learned how to manage environments, version data, and track experiments effectively. It has made my workflows more reliable, organized, and easier to share. I now feel confident setting up experiments that others can replicate and build upon. Reproducibility is not just good practice, it's essential for collaboration, scalability, and long-term impact in any ML project.